

2020 春

LINUX OPERATING SYSTEM

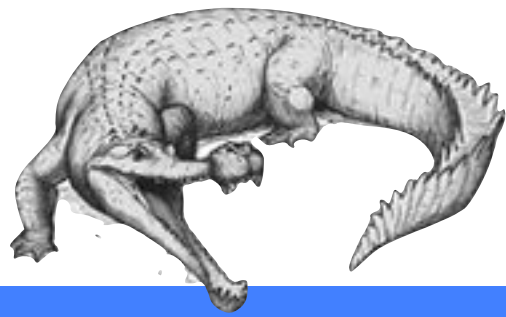
YANG

LINUX操作系统（双语）





双语课→课件内容中英混排



|Lecture 14

Page Table

本讲内容

 页表

 快表

 基于页的保护与共享

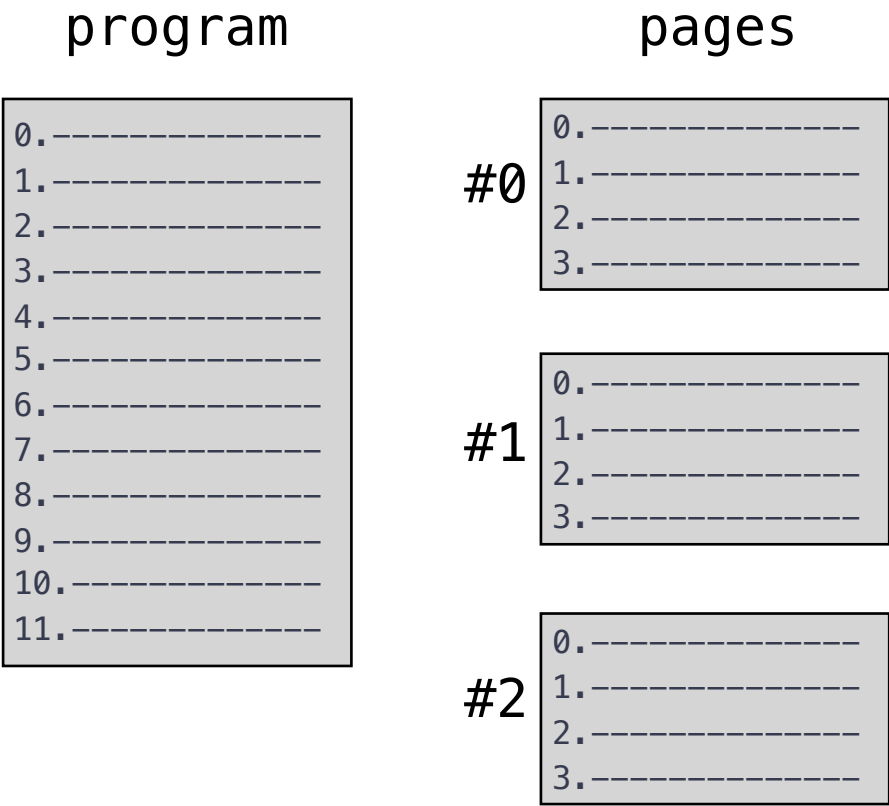
 多级页表

页表

分页

Frame Table

frame no.	state
0	0
1	0
2	0
3	0
4	0
5	0



页号	页内位移
----	------

逻辑地址

frame#0

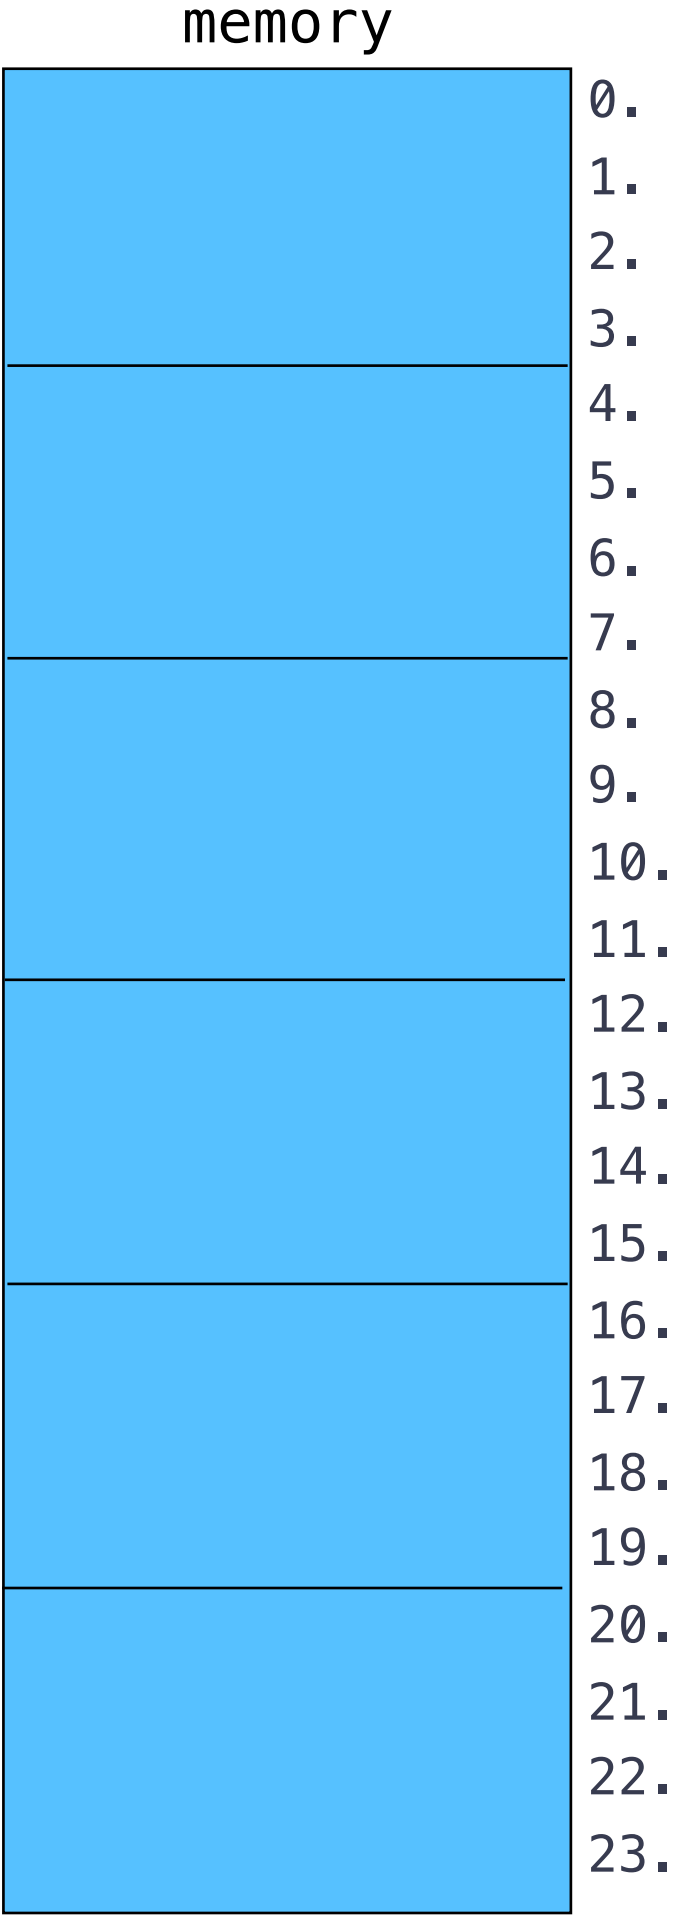
frame#1

frame#2

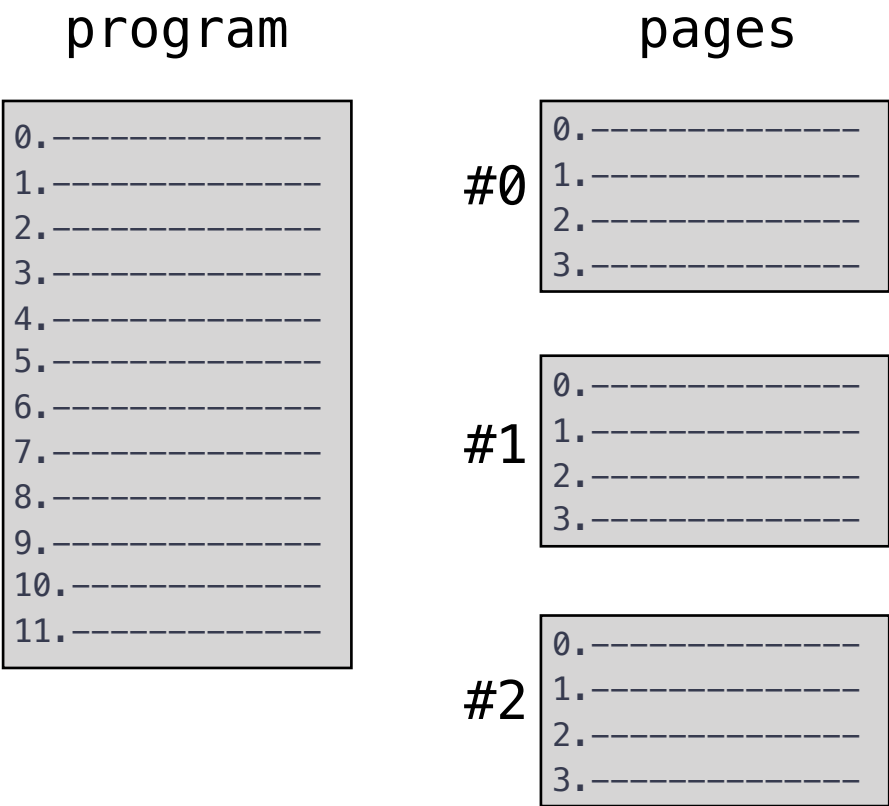
frame#3

frame#4

frame#5



分页



Frame Table

frame no.	state
0	1
1	0
2	1
3	0
4	0
5	1

Page Table

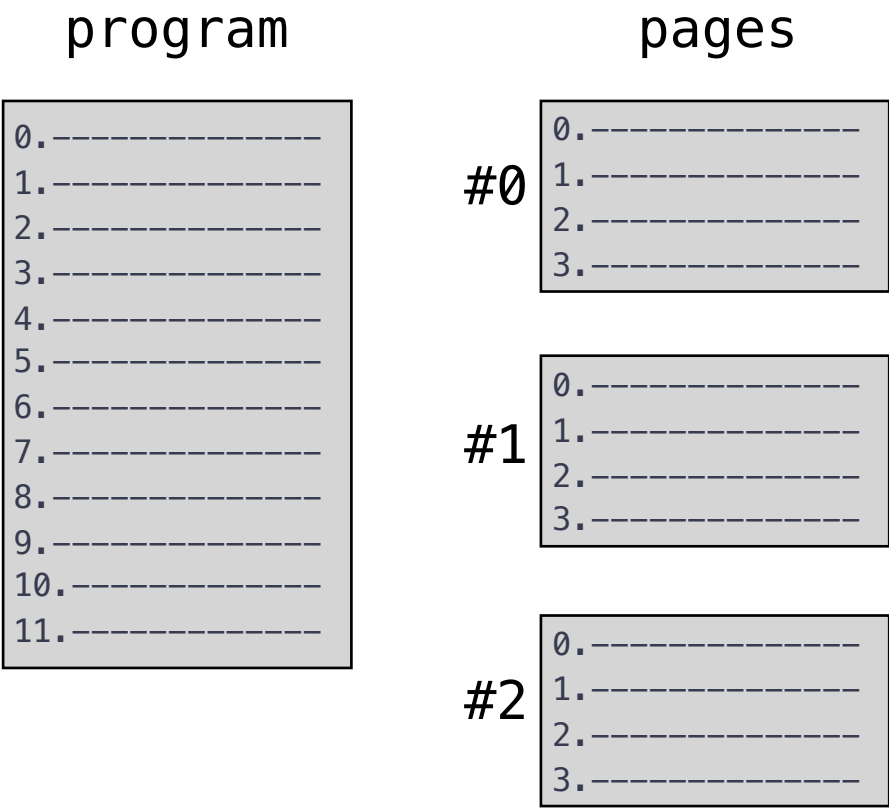
page no.	frame no.
0	5
1	0
2	2

页号	页内位移
----	------

逻辑地址



分页



Frame Table

frame no.	state
0	1
1	0
2	1
3	0
4	0
5	1

Page Table

	frame no.
0	5
1	0
2	2

页号	页内位移
----	------

逻辑地址



页面大小



逻辑地址

💡 若逻辑地址长度为 m bits，页面大小： 2^n Bytes

💡 页内位移占 n bits

💡 页号占 $m-n$ bits

💡 获取Linux系统页大小的命令

```
youngyt@youngyt-PC:/$ uname -m
x86_64
youngyt@youngyt-PC:/$ getconf PAGESIZE
4096
```

💡 请分别给出 m 和 n 的值

PAGE TABLE

- 💡 The operating system maintains a copy of the **page table** for **each process**.
- 💡 This copy is used to **translate logical addresses to physical addresses**.
- 💡 It is also used by the CPU dispatcher to define the **hardware page table** when a process is to be allocated the CPU.
- 💡 Paging therefore **increases the context-switch time**.

快表

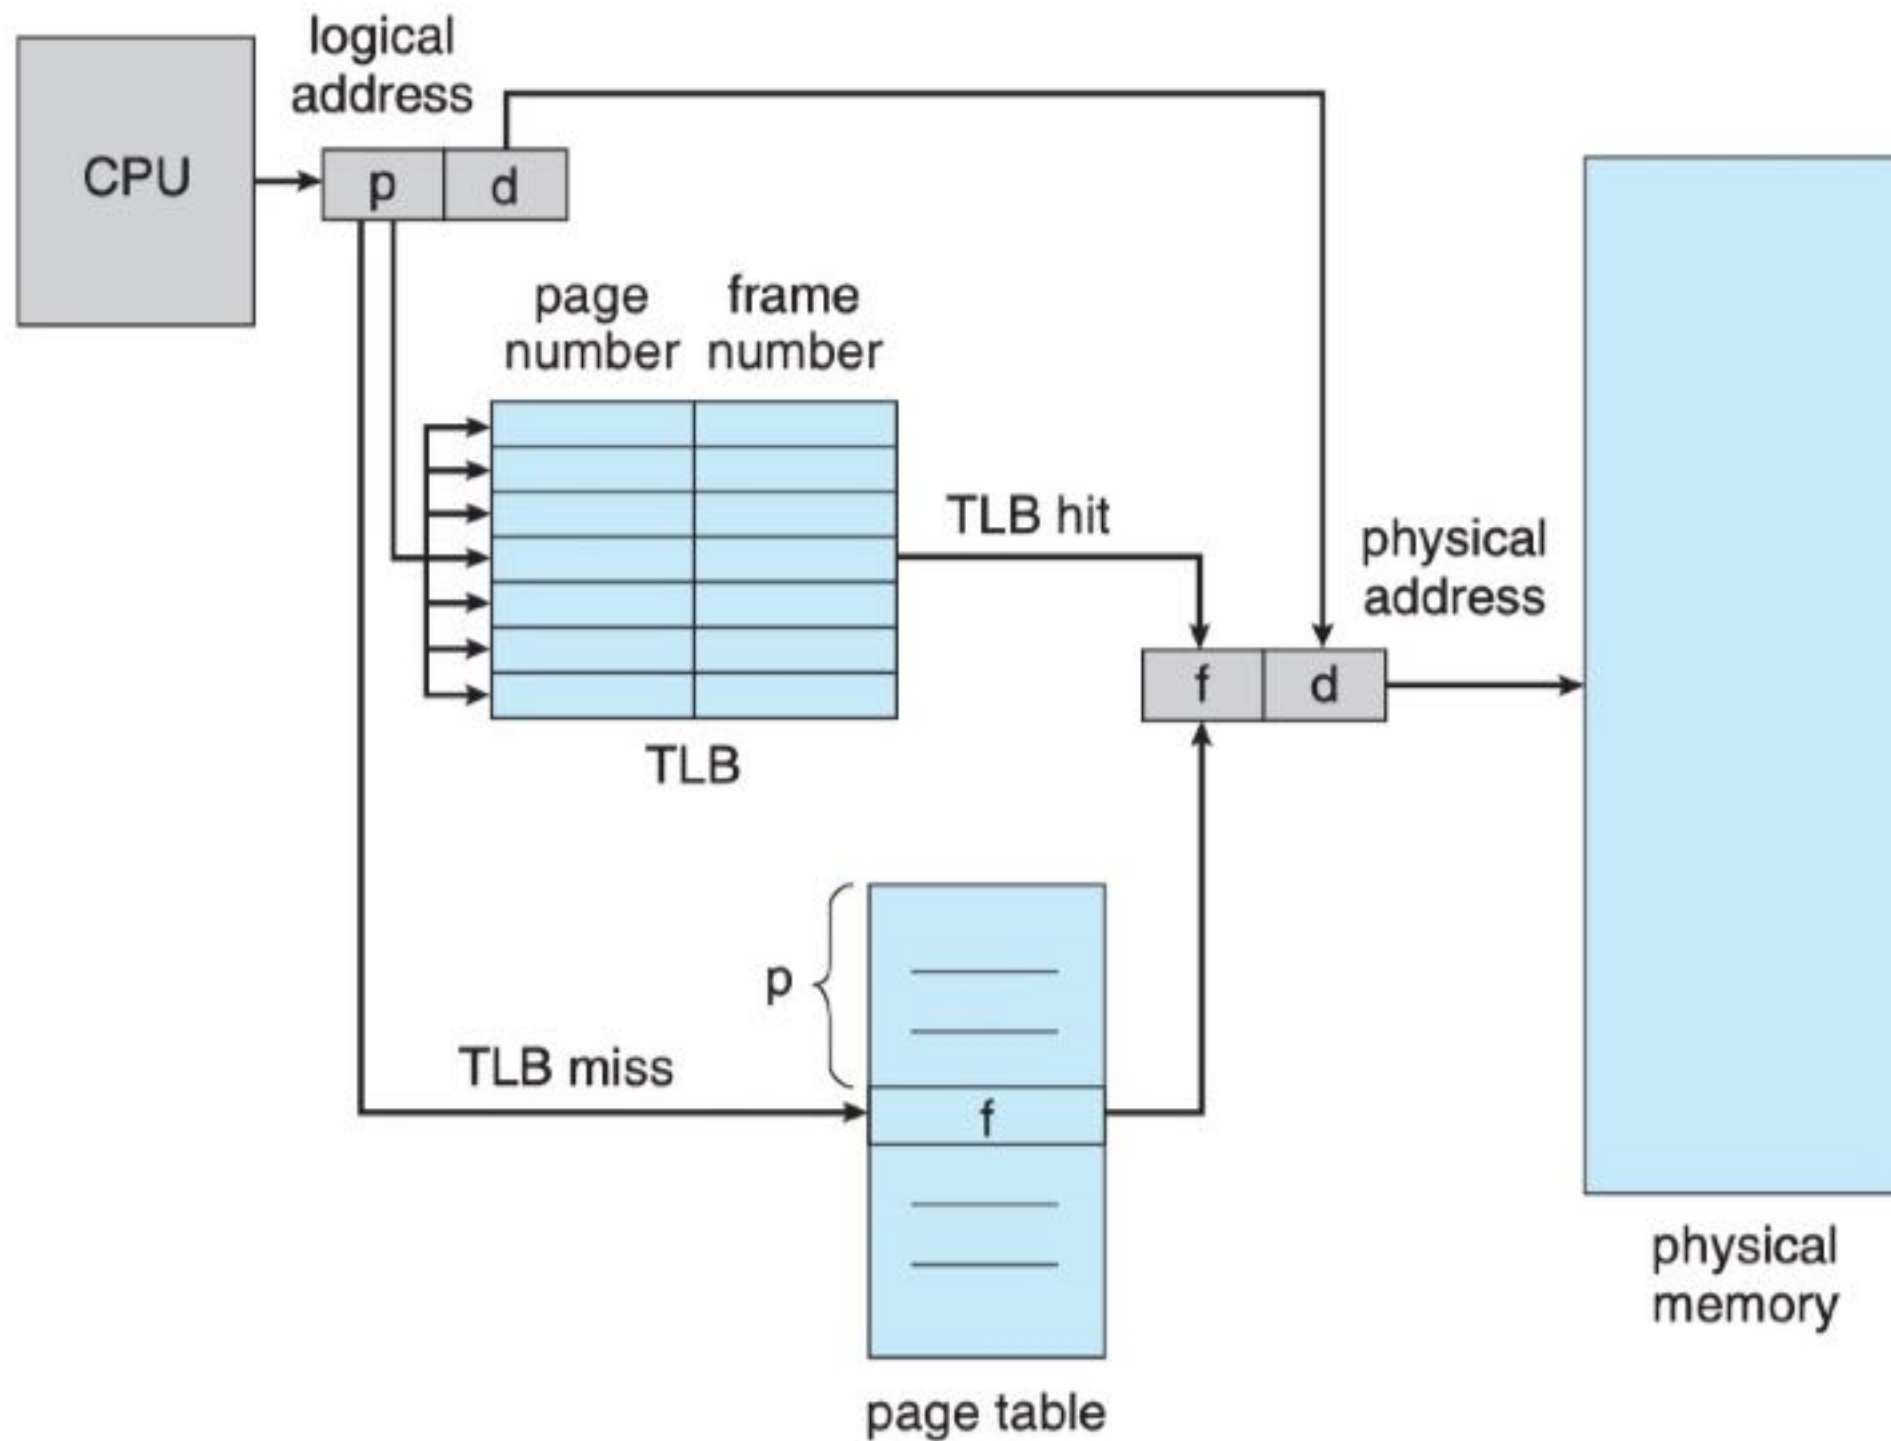
HARDWARE PAGE TABLE

- 💡 The page table is kept in main memory, and a page-table base register (PTBR) points to the page table.
- 💡 Changing page tables requires changing only this one register, substantially reducing context-switch time.
- 💡 With this scheme, two memory accesses are needed to access a byte (one for the page-table entry, one for the byte).

TLB

- 💡 TLB(Translation Look-aside Buffer) is a kind of small, fast-lookup hardware cache. It is used with page tables in the following way.
- 💡 The TLB contains only a few of the page-table entries.
- 💡 When a logical address is generated by the CPU, its page number is presented to the TLB.
- 💡 If the page number is found, its frame number is immediately available and is used to access memory.
- 💡 If TLB miss, a memory reference to the page table must be made.

PAGING WITH TLB



TLB HIT RATIO

💡 The percentage of times that the page number of interest is found in the TLB is called the **hit ratio**.

💡 An **80-percent hit ratio**, for example, means that we find the desired page number in the TLB 80 percent of the time. If it takes **100 nanoseconds** to access memory, please find the **effective memory-access time**.

$$\text{effective access time} = 0.80 \times 100 + 0.20 \times 200 = 120 \text{ ns}$$

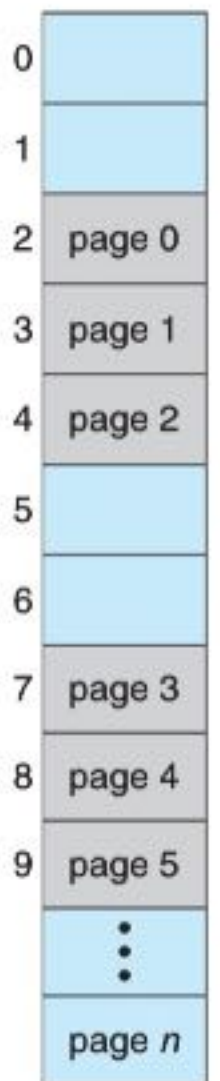
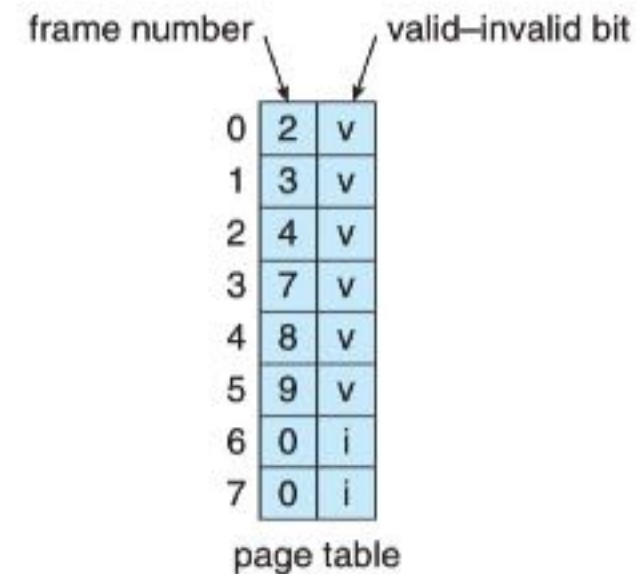
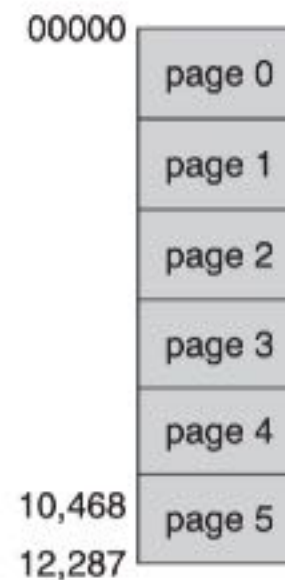
💡 How about 99% hit ratio?

基于页的保护与共享

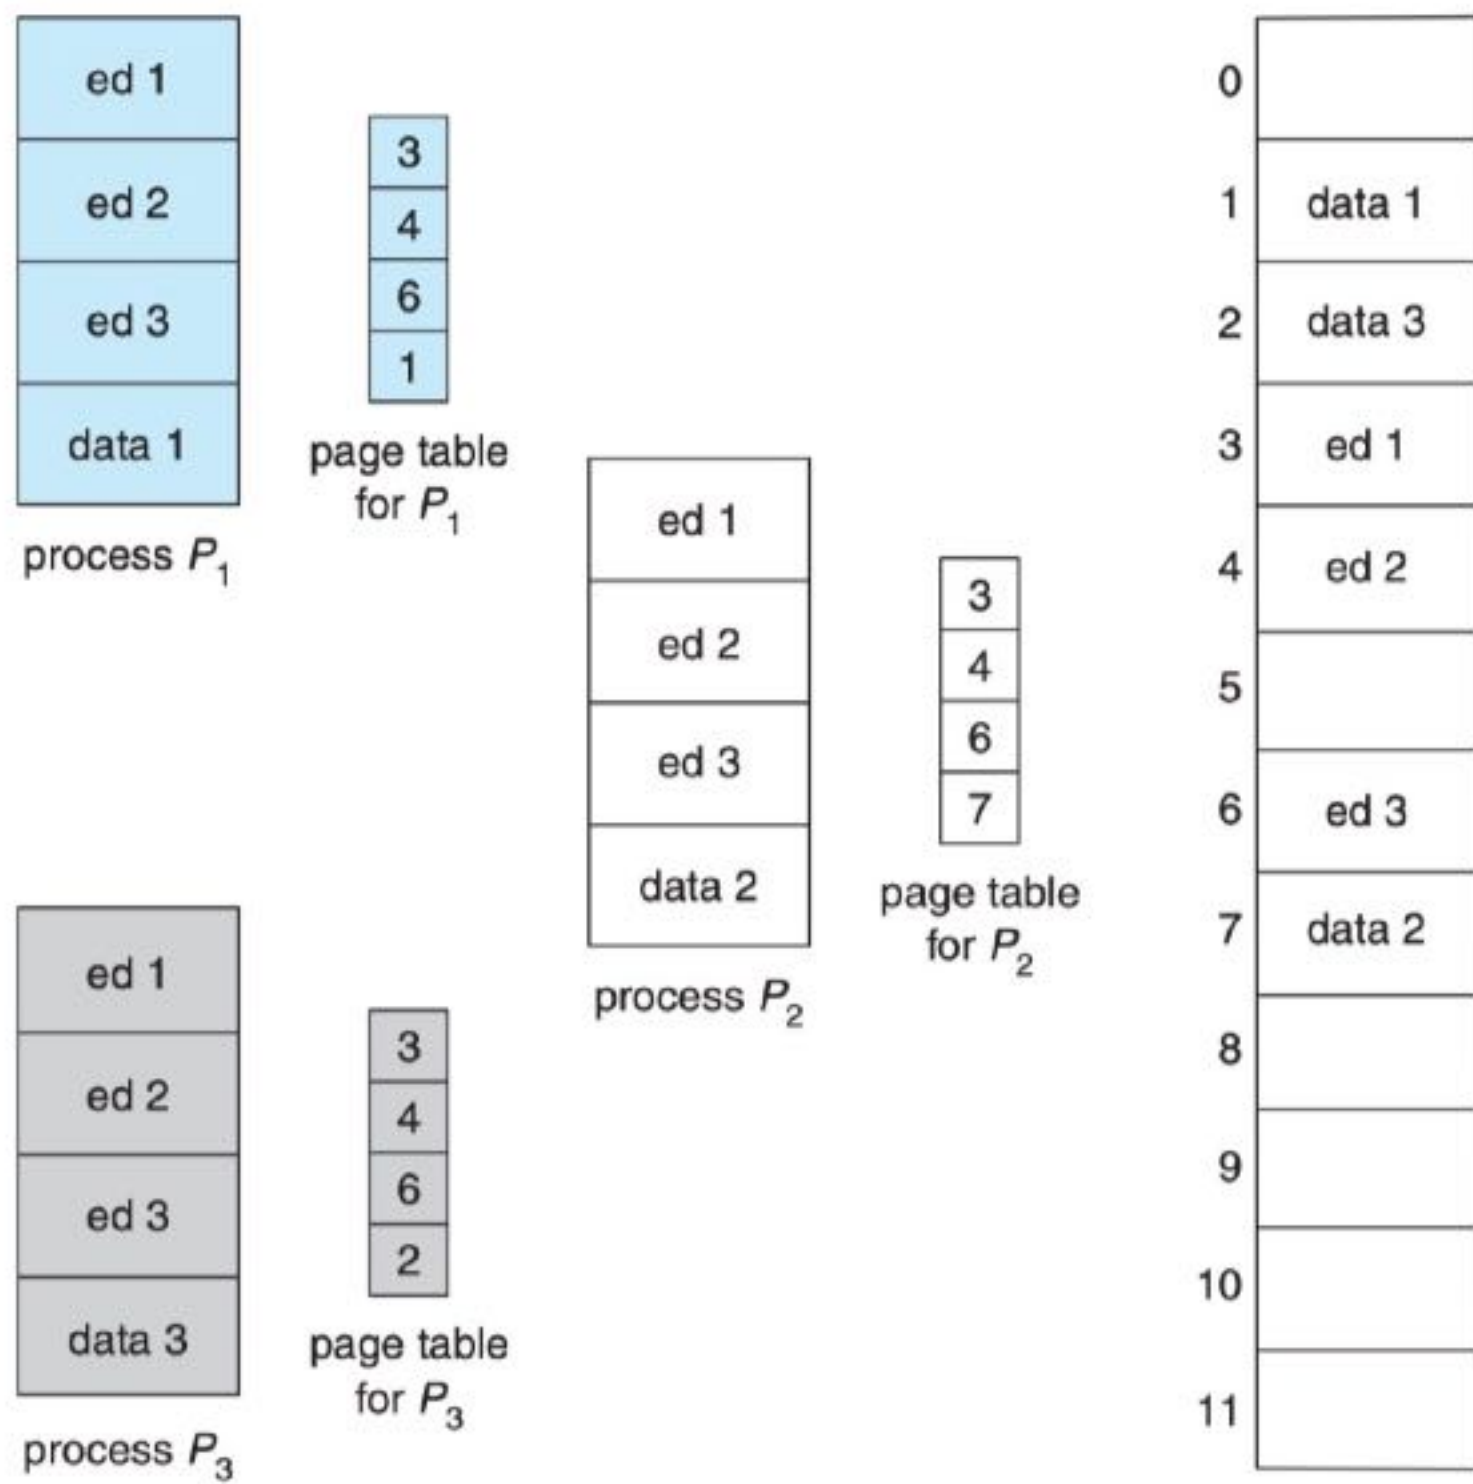
保护

💡 为了防止地址转换时出现异常，可在页表每个条目设置一个“**valid-invalid**”比特位，用于表示该页的有效性。

💡 这个方法可以被轻松扩展以提供更好的保护级别，如“只读”、“读写”、“可执行”等。



共享



多级页表

页表大小

- 💡 假设CPU是32bits，采用的逻辑地址是32bits，那么进程的逻辑地址空间大小为 2^{32} Bytes，即4G Bytes。
- 💡 若页面大小是4K Bytes，则一个进程最多被分成1M个页面，也就是说进程的页表最多有1M个页表项；
- 💡 若每个页表项占用4Bytes，则每个页表最多占用4MBytes空间（1K个连续页框）。
- 💡 如何解决“连续”的困扰？

页表页

pages

P#0
P#1
P#2
P#3
P#4
P#5
P#6
P#7

frames

F#0
F#1
F#2
F#3
F#4
F#5
F#6
F#7
F#8
F#9
F#10
F#11
F#12
F#13
F#14
F#15
F#16
F#17
F#18
F#19
F#20
F#21
F#22
F#23

页表页

pages

P#0
P#1
P#2
P#3
P#4
P#5
P#6
P#7

page table

P#0	F#3
P#1	F#6
P#2	F#12
P#3	F#9
P#4	F#1
P#5	F#8
P#6	F#11
P#7	F#0

frames

F#0
F#1
F#2
F#3
F#4
F#5
F#6
F#7
F#8
F#9
F#10
F#11
F#12
F#13
F#14
F#15
F#16
F#17
F#18
F#19
F#20
F#21
F#22
F#23

页表页

pages

P#0
P#1
P#2
P#3
P#4
P#5
P#6
P#7

page table

P#0	F#3
P#1	F#6
P#2	F#12
P#3	F#9
P#4	F#1
P#5	F#8
P#6	F#11
P#7	F#0

page table
directory

PPT#0	F#7
PPT#1	F#10

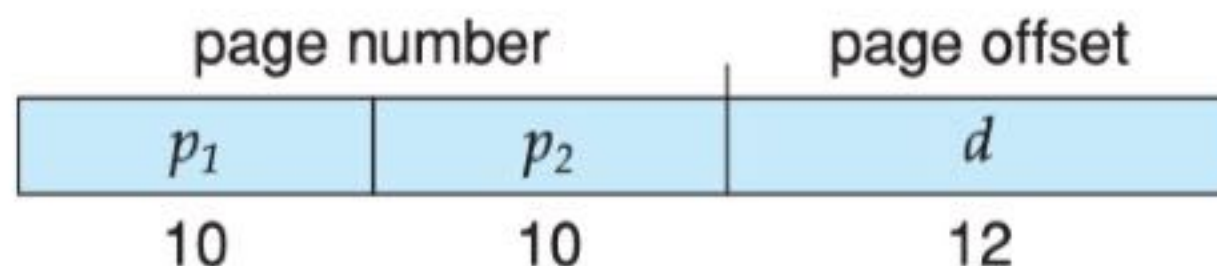
frames

F#0
F#1
F#2
F#3
F#4
F#5
F#6
F#7
F#8
F#9
F#10
F#11
F#12
F#13
F#14
F#15
F#16
F#17
F#18
F#19
F#20
F#21
F#22
F#23

页表页号	页号	页内位移
------	----	------

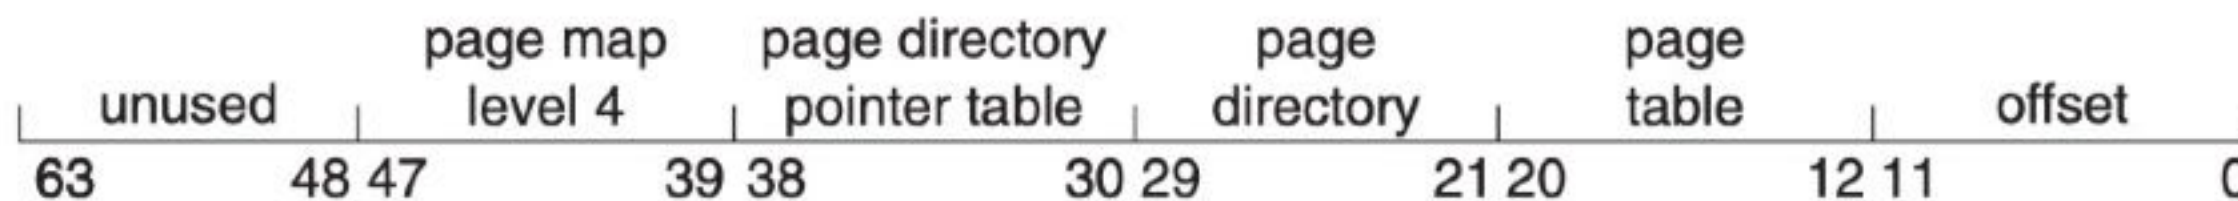
逻辑地址

多级页表



上面是一个32位地址采用**两级页表**的例子，页面大小是4KBytes，第一级页表页的数量是1K个，每个页表页中包含的页面数量也是1K个。

下面是x86-64架构CPU采用的**四级页表**方案：



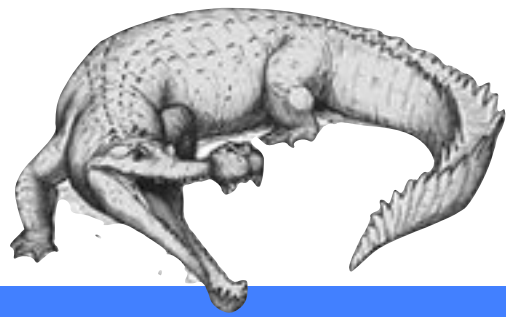
下期预告

 下次直播时间：3月25日 上午9:30

 课程内容

 Lecture 15 Virtual Memory

 Q&A



|Lecture 14

The End